

Uncovering Fraud Networks in a Blockchain Payment System

Background

This notebook analyzes the Elliptic Bitcoin Transaction Dataset to investigate potentially suspicious transaction networks in a blockchain payment system. The dataset contains transaction-level features, class labels, and directed payment flows, allowing us to study how funds move through the network, identify structurally important nodes, and detect patterns that may indicate illicit activity such as layering or fan-out behavior. The goal is to prepare the data carefully, build a meaningful subgraph for analysis, and produce evidence-based findings that are suitable for a compliance review.

1. Data Preparation & Validation

```
In [45]: # Import relevant packages
import pandas as pd
import numpy as np
```

```
In [46]: # Define file paths (adjust if your folder structure is different)
features_path = "elliptic_bitcoin_dataset/elliptic_txs_features.csv"
classes_path = "elliptic_bitcoin_dataset/elliptic_txs_classes.csv"
edges_path = "elliptic_bitcoin_dataset/elliptic_txs_edgelist.csv"
```

```
In [47]: # Load the datasets.

features = pd.read_csv(features_path, header=None)
# The features file, as previewed from Excel, doesn't have headers in its raw form.
# Column 0 is txId, Column 1 is time_step, and the rest are anonymized features.
# Rename the first two columns for clarity
features.rename(columns={0: 'txId', 1: 'time_step'}, inplace=True)

classes = pd.read_csv(classes_path)
edges = pd.read_csv(edges_path)
```

```
In [48]: # Check the first few rows
print("\nFeatures preview:")
print(features.head())

print("\nClasses preview:")
print(classes.head())

print("\nEdges preview:")
print(edges.head())
```

Features preview:

	txId	time_step	2	3	4	5	6	\
0	230425980	1	-0.171469	-0.184668	-1.201369	-0.121970	-0.043875	
1	5530458	1	-0.171484	-0.184668	-1.201369	-0.121970	-0.043875	
2	232022460	1	-0.172107	-0.184668	-1.201369	-0.121970	-0.043875	
3	232438397	1	0.163054	1.963790	-0.646376	12.409294	-0.063725	
4	230460314	1	1.011523	-0.081127	-1.201369	1.153668	0.333276	

	7	8	9	...	157	158	159	160	\
0	-0.113002	-0.061584	-0.162097	...	-0.562153	-0.600999	1.461330	1.461369	
1	-0.113002	-0.061584	-0.162112	...	0.947382	0.673103	-0.979074	-0.978556	
2	-0.113002	-0.061584	-0.162749	...	0.670883	0.439728	-0.979074	-0.978556	
3	9.782742	12.414558	-0.163645	...	-0.577099	-0.613614	0.241128	0.241406	
4	1.312656	-0.061584	-0.163523	...	-0.511871	-0.400422	0.517257	0.579382	

	161	162	163	164	165	166
0	0.018279	-0.087490	-0.131155	-0.097524	-0.120613	-0.119792
1	0.018279	-0.087490	-0.131155	-0.097524	-0.120613	-0.119792
2	-0.098889	-0.106715	-0.131155	-0.183671	-0.120613	-0.119792
3	1.072793	0.085530	-0.131155	0.677799	-0.120613	-0.119792
4	0.018279	0.277775	0.326394	1.293750	0.178136	0.179117

[5 rows x 167 columns]

Classes preview:

	txId	class
0	230425980	unknown
1	5530458	unknown
2	232022460	unknown
3	232438397	2
4	230460314	unknown

Edges preview:

	txId1	txId2
0	230425980	5530458
1	232022460	232438397
2	230460314	230459870
3	230333930	230595899
4	232013274	232029206

```
In [49]: # Check the shape (number of rows and columns) of each file
print("Features shape:", features.shape)
print("Classes shape:", classes.shape)
print("Edges shape:", edges.shape)
```

Features shape: (203769, 167)

Classes shape: (203769, 2)

Edges shape: (234355, 2)

```
In [50]: # Check for duplicate transaction IDs in features and classes
print("Duplicate txId in features:", features["txId"].duplicated().sum())
print("Duplicate txId in classes:", classes["txId"].duplicated().sum())

# Check for duplicate edges
print("Duplicate edges:", edges.duplicated().sum())
```

Duplicate txId in features: 0

Duplicate txId in classes: 0

Duplicate edges: 0

All three files are structurally clean with no duplicate transaction IDs or repeated edges, confirming data integrity ahead of merging and graph construction.

```
In [51]: # Check missing values
print("\nMissing values in features:")
print(features.isna().sum().sum())

print("\nMissing values in classes:")
print(classes.isna().sum())

print("\nMissing values in edges:")
print(edges.isna().sum())
```

Missing values in features:
0

Missing values in classes:
txId 0
class 0
dtype: int64

Missing values in edges:
txId1 0
txId2 0
dtype: int64

There are no missing values in any of the dataset files

```
In [52]: # Check the distribution in the classes file
print("\nClass distribution:")
print(classes["class"].value_counts(dropna=False))
```

Class distribution:
class
unknown 157205
2 42019
1 4545
Name: count, dtype: int64

The class labels are highly imbalanced, with most transactions unlabelled, far fewer licit nodes, and only a small illicit subset.

```
In [53]: # Merge features with classes using txID
merged_df = features.merge(classes, on="txId", how="left")

# Check the merged result
print("Merged shape:", merged_df.shape)
```

Merged shape: (203769, 168)

```
In [54]: # Check if any class labels are still missing after merge
print("\nMissing class labels after merge:")
print(merged_df["class"].isna().sum())
```

Missing class labels after merge:
0

```
In [55]: # Extract all unique transaction IDs that appear in the edge flows
unique_edge_nodes = set(edges['txId1']).union(set(edges['txId2']))
unique_master_nodes = set(merged_df['txId'])

# Find if any transactions in the edge list are missing from our master node records
orphan_nodes = unique_edge_nodes - unique_master_nodes

print(f"Total unique nodes involved in flows: {len(unique_edge_nodes)}")
print(f"Orphan transactions (in edge list but missing from master records): {len(orphan_nodes)}
```

Total unique nodes involved in flows: 203769

Orphan transactions (in edge list but missing from master records): 0

All transactions involved in the edge list are already present in the master node table (merged_df), so the graph is structurally complete.

Justification

- I kept records with "unknown" labels in classes, because unlabeled transactions are still part of the real network and may later appear suspicious through their position or behavior.
- I used left join to merge features and classes because it keeps every transaction in the features table (main table) and only adds labels where a match exists, which preserves the full network for analysis.
- features and classes tables were merged into a master node dataframe (merged_df) via txId; the edges file was intentionally excluded from a tabular join and instead integrated structurally via NetworkX as a Directed Graph, since a row-wise merge would artificially replicate transaction rows and corrupt node-level features, violating data integrity standards required for regulatory compliance (NFIU/CBN).

2. Network Construction & Subgraph Selection

```
In [56]: # Import relevant package
import networkx as nx
import random
```

a. Build the Graph

```
In [57]: # Create a full directed graph
G_full = nx.DiGraph()
```

```
In [58]: # Add edges from the edge list
# txId1 = source transaction
# txId2 = destination transaction
G_full.add_edges_from(edges[['txId1', 'txId2']].itertuples(index=False, name=None))
```

```
In [59]: # Check basic graph size
print("Number of nodes:", G_full.number_of_nodes())
print("Number of edges:", G_full.number_of_edges())
```

Number of nodes: 203769
Number of edges: 234355

The graph is complete and fully connected to the dataset scale, with 203,769 nodes and 234,355 edges, so the edge list has been successfully captured as the full transaction network.

b. Add node attributes

Attach the class labels and time step from your merged node table so later analysis can color nodes and detect suspicious patterns.

```
In [60]: # Create dictionaries for node attributes
class_dict = merged_df.set_index("txId")["class"].to_dict()
time_dict = merged_df.set_index("txId")["time_step"].to_dict()

# The dataframe was converted to dictionaries of records for fast loading
```

```
In [61]: # Add attributes to graph nodes
nx.set_node_attributes(G_full, class_dict, "class")
nx.set_node_attributes(G_full, time_dict, "time_step")
```

```
In [62]: # Check one sample node
sample_node = list(G_full.nodes())[0]
print(sample_node, G_full.nodes[sample_node])

230425980 {'class': 'unknown', 'time_step': 1}
```

c. Random sampling for betweenness

Compute betweenness centrality on a random sample of 20,000 nodes, then select the top 50 most structurally important transactions.

```
In [63]: # Set seed for reproducibility
random.seed(4)
```

```
In [64]: # Take a random sample of 50,000 nodes
sample_nodes = random.sample(list(G_full.nodes()), 20000)
```

```
In [65]: # Create the sampled subgraph
G_sample = G_full.subgraph(sample_nodes).copy()

print("Sampled graph nodes:", G_sample.number_of_nodes())
print("Sampled graph edges:", G_sample.number_of_edges())
```

Sampled graph nodes: 20000
Sampled graph edges: 2246

d. Compute betweenness

Calculate betweenness on the sampled graph.

```
In [66]: # Compute betweenness centrality on the sampled graph
bet_sample = nx.betweenness_centrality(G_sample, normalized=True)
```

```
In [67]: # Convert to DataFrame and get top 50 nodes
bet_df = pd.DataFrame(
    bet_sample.items(),
    columns=["txId", "betweenness"]
).sort_values("betweenness", ascending=False)

top_50_nodes = bet_df.head(50)["txId"].tolist()

print(bet_df.head(10))
```

	txId	betweenness
1681	8923841	3.000450e-08
4725	22837965	2.000300e-08
19935	9829997	2.000300e-08
3406	8935133	1.500225e-08
6031	244751101	1.500225e-08
649	209064416	1.250188e-08
4050	244738146	1.000150e-08
7074	194950104	1.000150e-08
2878	244730003	1.000150e-08
16484	11511112	1.000150e-08

e. Extract ego network

```
In [68]: # Build ego networks around the top 50 nodes
ego_graphs = [nx.ego_graph(G_full, node, radius=3, undirected=True) for node in top_50]
```

For directed graphs, `ego_graph()` by default follows outgoing neighbors, but using `undirected=True` helps include both in- and out-neighbours when building a fuller working subgraph.

```
In [69]: # Combine them into one working subgraph
working_subgraph = nx.compose_all(ego_graphs)

print("Working subgraph nodes:", working_subgraph.number_of_nodes())
print("Working subgraph edges:", working_subgraph.number_of_edges())
```

```
Working subgraph nodes: 7632
Working subgraph edges: 12703
```

Justification

Expanding the subgraph extraction to a three hop radius (`radius=3`) was a necessary technical pivot to meet the 5,000 node requirement while navigating severe computational constraints. Initial attempts to compute approximated centrality across the entire network exceeded processing limits. To bypass this, we isolated a random structural sample of 20,000 nodes and calculated exact centrality within that subset to identify localized hubs. Because this isolated random sampling inherently fragments the broader network topology, extracting only their immediate neighbors yielded an insufficient node count.

The primary analytical advantage of increasing the extraction depth to three hops is that it not only reconnects these fragmented communities but also maps the extended lifecycle of complex transactions. This deeper visibility successfully captures the multi step obfuscation and

layering tactics typical of illicit blockchain activity, delivering a highly comprehensive subgraph strictly on deadline.

f. Computing metrics for the working subgraph

```
In [70]: # Calculate metrics
in_degree = dict(working_subgraph.in_degree())
out_degree = dict(working_subgraph.out_degree())
sub_betweenness = nx.betweenness_centrality(working_subgraph)
pagerank = nx.pagerank(working_subgraph)
```

```
In [71]: # Store metrics in a DataFrame
metrics_df = pd.DataFrame({
    'txId': list(working_subgraph.nodes()),
    'in_degree': [in_degree[n] for n in working_subgraph.nodes()],
    'out_degree': [out_degree[n] for n in working_subgraph.nodes()],
    'betweenness': [sub_betweenness[n] for n in working_subgraph.nodes()],
    'pagerank': [pagerank[n] for n in working_subgraph.nodes()]
})
```

```
In [72]: # Extract class labels from your master merged_df for the nodes in our subgraph
subgraph_classes = merged_df[merged_df['txId'].isin(working_subgraph.nodes())][['txId', 'class']]

# Merge classes into the metrics DataFrame
metrics_df = metrics_df.merge(subgraph_classes, on='txId', how='left')

# Identify the top 10 most central nodes based on PageRank
top_10_nodes = metrics_df.nlargest(10, 'pagerank')['txId'].tolist()

print("\nMetrics computed successfully. Top 10 Hubs identified.")
print(metrics_df.head())
```

Metrics computed successfully. Top 10 Hubs identified.

	txId	in_degree	out_degree	betweenness	pagerank	class
0	8927237	3	0	0.000000e+00	0.000202	unknown
1	8927242	5	8	2.937024e-06	0.000182	unknown
2	121854988	1	0	0.000000e+00	0.000106	unknown
3	8860181	2	1	2.404485e-08	0.000106	unknown
4	8923183	9	6	3.776188e-06	0.000342	unknown

g. Network Visualization

```
In [73]: # Import relevant packages
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
```

```
In [79]: # Prepare visual attributes
color_map = []
node_sizes = []

for node in working_subgraph.nodes():
    # Fetch the class for the current node
    node_class_series = metrics_df.loc[metrics_df['txId'] == node, 'class'].values
    node_class = node_class_series[0] if len(node_class_series) > 0 else 'unknown'

    # Sizing: Make the top 10 nodes much larger
    if node in top_10_nodes:
```

```

        node_sizes.append(300)
    else:
        node_sizes.append(25)

    # Coloring: Distinct colors for classes and highlighting top 10
    if node in top_10_nodes:
        color_map.append('gold')
    elif node_class == '1': # Illicit
        color_map.append('red')
    elif node_class == '2': # Licit
        color_map.append('green')
    else: # Unknown
        color_map.append('lightgray')

# Plot setup
plt.figure(figsize=(16, 12))
plt.title("Blockchain Subgraph: Network Topology and Key Hubs", fontsize=16)

# Generate Layout
# k=0.15 acts as the optimal spacing distance to spread out the nodes
print("Generating visualization layout...")
pos = nx.spring_layout(working_subgraph, k=0.01, seed=42)


# Draw the network
nx.draw(
    working_subgraph,
    pos,
    node_color=color_map,
    node_size=node_sizes,
    edge_color='whitesmoke',
    alpha=0.85,
    with_labels=False
)

# Custom Legend to map colors to definitions
legend_elements = [
    mpatches.Patch(color='red', label='Illicit Nodes (Class 1)'),
    mpatches.Patch(color='green', label='Licit Nodes (Class 2)'),
    mpatches.Patch(color='lightgray', label='Unknown Nodes'),
    mpatches.Patch(color='gold', label='Top 10 Central Hubs')
]
plt.legend(handles=legend_elements, loc='upper right', fontsize=12)

# Render plot
plt.tight_layout()
plt.show()

```

Generating visualization layout...



Inference

- Legitimate Central Hubs: The largest central nodes are surrounded by licit and unlabelled transactions, indicating that fraudsters route funds through major legitimate exchanges to blend in, rather than owning the hubs themselves.
- Peripheral Layering: Illicit (red) nodes are scattered along the outer edges of the network, visually confirming that bad actors use multiple hops to obfuscate the origin of funds before moving them toward the center.
- Unknown Intermediaries: Linear "spoke" patterns at the edges reveal fan-out behavior, where funds are rapidly broken down and moved through clusters of unlabelled (grey) nodes acting as temporary accomplice wallets.

3. Fraud Pattern Detection

Detect the Patterns Programmatically

```
In [80]: # 1. Bring in the time_step data from your master merged_df
node_times = merged_df[merged_df['txId'].isin(working_subgraph.nodes())][['txId', 'tin
```

```
# Ensure metrics_df has the latest data
task3_df = metrics_df.merge(node_times, on='txId', how='left')
```

```
In [81]: # 2. Define our typologies
# Fan-Out: Receives from 1 or 2 sources, but sends to 10+ distinct destinations
fan_out_flags = task3_df[(task3_df['in_degree'] <= 2) & (task3_df['out_degree'] >= 10)]
fan_out_flags['pattern'] = 'Fan-Out'

# Layering: Receives from multiple sources (3+) and quickly redistributes to multiple
layering_flags = task3_df[(task3_df['in_degree'] >= 3) & (task3_df['out_degree'] >= 3)]
layering_flags['pattern'] = 'Layering'

# Combine the flagged nodes
flagged_nodes = pd.concat([fan_out_flags, layering_flags]).drop_duplicates(subset=['tx
```

```
In [82]: # 3. Output the specific documentation metrics
print(f"\n--- FAN-OUT PATTERNS DETECTED ({len(fan_out_flags)} nodes) ---")
print(fan_out_flags[['txId', 'in_degree', 'out_degree', 'time_step', 'class']].head(10))

print(f"\n--- LAYERING BEHAVIOR DETECTED ({len(layering_flags)} nodes) ---")
print(layering_flags[['txId', 'in_degree', 'out_degree', 'time_step', 'class']].head(10))

--- FAN-OUT PATTERNS DETECTED (91 nodes) ---
   txId  in_degree  out_degree  time_step  class
8931887         1          11          6  unknown
8923819         2          12          6  unknown
8923825         2          13          6  unknown
8937835         0          13          6  unknown
3402358         0          12          4  unknown
1416324         0          14          4  unknown
64949778        1          10          4  unknown
2512477         0          22          4      2
244721166        0          10          3  unknown
244722777        2          14          3  unknown

--- LAYERING BEHAVIOR DETECTED (295 nodes) ---
   txId  in_degree  out_degree  time_step  class
8927242         5           8          6  unknown
8923183         9           6          6  unknown
121656897        5           5          6  unknown
8929881         8           8          6  unknown
8950367         7           4          6  unknown
8940130         5           5          6  unknown
9829997         8           8          6  unknown
8942708         3           7          6  unknown
8929408         8           7          6  unknown
8929409         4          10          6  unknown
```

```
In [83]: # 4. Cross-Reference with Class Labels
class_breakdown = flagged_nodes['class'].value_counts(dropna=False)
print("\n--- CLASS LABEL CROSS-REFERENCE ---")
print(f"Flagged nodes labelled Illicit (Class 1): {class_breakdown.get('1', 0)}")
print(f"Flagged nodes labelled Licit (Class 2): {class_breakdown.get('2', 0)}")
print(f"Flagged nodes labelled Unknown: {class_breakdown.get('unknown', 0)}")

--- CLASS LABEL CROSS-REFERENCE ---
Flagged nodes labelled Illicit (Class 1): 0
Flagged nodes labelled Licit (Class 2): 67
Flagged nodes labelled Unknown: 319
```

Fraud Pattern Detection Findings

Fan-Out Patterns

- What it is: A network topology where a node receives funds from a minimal number of sources and rapidly disperses them to a large number of distinct destination wallets within a single time step.
- Nodes Involved: A total of 91 nodes exhibited this behavior. For example, Node 8931887 received funds from 1 source and dispersed them to 11 unique destinations at time step 6. Node 2512477 sent funds to 22 distinct destinations at time step 4.
- Compliance Risk: This structure is a primary indicator of the placement phase of money laundering. Bad actors utilize fan-out transactions to break down large, easily detectable sums of cryptocurrency into micro-transactions, often spreading them across unhosted wallets to evade Anti-Money Laundering detection thresholds.

Layering Behavior

- What it is: Nodes exhibiting simultaneously high in-degree and high out-degree centrality within a condensed time frame, acting as high-throughput intermediaries.
- Nodes Involved: A total of 295 nodes exhibited this behavior. For example, Node 8923183 received connections from 9 sources and redistributed to 6 destinations at time step 6. Node 8927242 consolidated funds from 5 sources and sent them to 8 destinations at the exact same time step.
- Compliance Risk: This pattern strongly suggests the use of localized mixing services or peel chains. The node collects funds from multiple origins and immediately redistributes them, which is a deliberate obfuscation tactic designed to sever the on-chain link between the source and the final liquidation point.

Class Label Cross-Reference

- Flagged nodes labelled Illicit (Class 1): 0
- Flagged nodes labelled Licit (Class 2): 67
- Flagged nodes labelled Unknown: 319

Suspicion of Unlabelled Nodes: The absence of explicitly flagged Class 1 nodes in this subset does not mitigate the risk. The fact that 319 completely unlabelled nodes are executing the vast majority of the detected layering and fan-out patterns is highly suspicious. Their network position strongly implies these are temporary, intermediary accomplice wallets generated specifically to facilitate the movement of undocumented funds before they merge with the 67 licit nodes (likely exchanges).

Interactive HTML Dashboard

```
In [99]: import plotly.graph_objects as go
import plotly.express as px

print("Building Interactive Dashboard Components...")

# 1. Prepare and Map Data
# Map the numeric/string classes to intuitive labels
label_map = {'1': 'Illicit', '2': 'Licit', '0': 'Unknown', 'unknown': 'Unknown'}
task3_df['class_label'] = task3_df['class'].astype(str).map(label_map)

# Define the Composite Risk Score
task3_df['norm_betweenness'] = task3_df['betweenness'] / task3_df['betweenness'].max()
task3_df['norm_pagerank'] = task3_df['pagerank'] / task3_df['pagerank'].max()
task3_df['risk_score'] = (task3_df['norm_betweenness'] * 0.5) + (task3_df['norm_pagerank'] * 0.5)

# Sort for the top 20 riskiest nodes
top_20_nodes = task3_df.nlargest(20, 'risk_score')
top_10_ids = task3_df.nlargest(10, 'risk_score')['txId'].tolist()

# 2. Build the Risk Ranking Table
table_fig = go.Figure(data=[go.Table(
    header=dict(
        values=['Transaction ID', 'Class', 'In-Degree', 'Out-Degree', 'Risk Score'],
        fill_color='midnightblue',
        font=dict(color='white', size=12),
        align='left'
    ),
    cells=dict(
        values=[
            top_20_nodes['txId'],
            top_20_nodes['class_label'],
            top_20_nodes['in_degree'],
            top_20_nodes['out_degree'],
            top_20_nodes['risk_score'].round(4)
        ],
        fill_color='aliceblue',
        align='left'
    )
)])
table_fig.update_layout(
    title="Top 20 Riskiest Nodes",
    margin=dict(l=0, r=0, t=40, b=0),
    height=450
)

# 3. Build the Time-Based Chart
time_data = task3_df.groupby(['time_step', 'class_label']).size().reset_index(name='count')
```

```

time_fig = px.bar(
    time_data,
    x='time_step',
    y='count',
    color='class_label',
    color_discrete_map={'Illicit': 'red', 'Licit': 'green', 'Unknown': 'lightgray'},
    title="Transaction Volume by Time Step",
    labels={'time_step': 'Time Step', 'count': 'Number of Transactions', 'class_label':
})

time_fig.update_layout(
    margin=dict(l=0, r=0, t=40, b=0),
    height=450,
    showlegend=False
)

# 4. Build the Interactive Network Graph
edge_x = []
edge_y = []
for edge in working_subgraph.edges():
    x0, y0 = pos[edge[0]]
    x1, y1 = pos[edge[1]]
    edge_x.extend([x0, x1, None])
    edge_y.extend([y0, y1, None])

edge_trace = go.Scatter(
    x=edge_x,
    y=edge_y,
    line=dict(width=0.5, color='#ccc'),
    hoverinfo='none',
    mode='lines',
    showlegend=False
)

# Split nodes by class so the Legend can filter them natively
traces = [edge_trace]
categories = ['Illicit', 'Licit', 'Unknown']
color_mapping = {'Illicit': 'red', 'Licit': 'green', 'Unknown': 'lightgray'}

for cat in categories:
    cat_x, cat_y, cat_color, cat_size, cat_text = [], [], [], [], []
    for node in working_subgraph.nodes():
        node_info = task3_df[task3_df['txId'] == node].iloc[0]
        if node_info['class_label'] == cat:
            x, y = pos[node]
            cat_x.append(x)
            cat_y.append(y)
            cat_text.append(f"ID: {node}<br>Class: {cat}<br>Risk: {node_info['risk_score']}")

            if node in top_10_ids:
                cat_size.append(25)
                cat_color.append('gold')
            else:
                cat_size.append(10 if cat == 'Illicit' else (8 if cat == 'Licit' else 5))
                cat_color.append(color_mapping[cat])

    traces.append(go.Scatter(
        x=cat_x,
        y=cat_y,
        mode='markers',

```

```

        hoverinfo='text',
        name=cat,
        marker=dict(
            color=cat_color,
            size=cat_size,
            line_width=1,
            line_color='black'
        ),
        text=cat_text
    ))

net_fig = go.Figure(
    data=traces,
    layout=go.Layout(
        title=dict(
            text='Interactive Blockchain Network (Top 10 Hubs in Gold)',
            font=dict(size=20)
        ),
        hovermode='closest',
        margin=dict(b=0, l=0, r=0, t=40),
        xaxis=dict(showgrid=False, zeroline=False, showticklabels=False),
        yaxis=dict(showgrid=False, zeroline=False, showticklabels=False),
        height=720,
        showlegend=True,
        legend=dict(
            title=dict(text="Filter Network", font=dict(size=30, color="black")),
            font=dict(size=24, color="black"),
            bgcolor="rgba(255, 255, 255, 0.98)",
            bordercolor="black",
            borderwidth=2,
            yanchor="top",
            y=0.99,
            xanchor="right",
            x=0.995,
            itemsizing="constant",
            itemwidth=80,
            tracegroupgap=18,
            itemclick="toggle",
            itemdoubleclick="toggleothers"
        )
    )
)

# 5. Export to a Single HTML File
print("Compiling standalone HTML dashboard...")

html_content = f"""
<!DOCTYPE html>
<html>
<head>
    <title>Blockchain Fraud Network Dashboard</title>
    <script src="https://cdn.plot.ly/plotly-latest.min.js"></script>
    <style>
        body {{ font-family: Arial, sans-serif; background-color: #f4f7f6; margin: 2vh
        .header {{ text-align: center; padding: 15px; background-color: #2c3e50; color
        .header h1 {{ margin: 0; font-size: 24px; }}
        .header p {{ margin: 5px 0 0 0; font-size: 14px; }}
        .grid-container {{ display: grid; grid-template-columns: 1fr 1fr; gap: 15px; m
        .full-width {{ grid-column: 1 / span 2; background: white; padding: 10px; bord
        .half-width {{ background: white; padding: 10px; border-radius: 8px; box-shad

```

```

        .filter-note {{ text-align: center; color: #7f8c8d; font-style: italic; margin: 10px 0; }}
    </style>
</head>
<body>
    <div class="header">
        <h1>Fraud Network Analysis Dashboard</h1>
        <p>Interactive exploration of flagged Bitcoin transactions and structural hubs</p>
    </div>

    <div class="filter-note">
        Use the slicer on the right side of the Network Graph to filter the nodes by c
    </div>

    <div class="grid-container">
        <div class="full-width">
            {net_fig.to_html(full_html=False, include_plotlyjs=False)}
        </div>
        <div class="half-width">
            {table_fig.to_html(full_html=False, include_plotlyjs=False)}
        </div>
        <div class="half-width">
            {time_fig.to_html(full_html=False, include_plotlyjs=False)}
        </div>
    </div>
</body>
</html>
"""

with open("compliance_dashboard.html", "w", encoding="utf-8") as f:
    f.write(html_content)

print("SUCCESS: compliance_dashboard.html has been generated.")

```

Building Interactive Dashboard Components...

Compiling standalone HTML dashboard...

SUCCESS: compliance_dashboard.html has been generated.